

FedTS Enhancements: Federated Learning Simulation Report

Goal

The goal of this script is to **simulate a dynamic and robust Federated Learning (FL)** environment that supports:

- Comparison between **Thompson Sampling (Fed-TS)** and **Random client selection**
 - Integration of **malicious clients** (via label-flipping attacks)
 - Use of complex datasets like **CIFAR-10**
 - **Tracking and visualizing** model performance over rounds
 - Enhancing trust-based **client participation** in FL systems
 - Optional **model checkpoint saving** for reproducibility
-

Problem Statement

Traditional Federated Learning assumes that all participating clients behave honestly and contribute fairly. However, in real-world scenarios:

- Some clients may behave **maliciously** (e.g., label-flipping attacks).
- Client reliability may vary across training rounds.
- Static client participation doesn't reflect practical deployments.

Hence, we simulate:

- **Dynamic client participation** with malicious clients
- **Trust-aware selection mechanism** using **Fed-TS**

- Compare it against naive **random client selection**
-

✓ What This Script Does

1. Dataset Preparation

- Loads **CIFAR-10** (or MNIST if toggled) with `torchvision`.
- Partitions dataset among `CLIENTS_TOTAL` clients non-IID style.
- Adds **label-flipping attackers** among dynamically inserted clients.

2. Client Simulation

- `CLIENTS_ORIGINAL` are the fixed set.
- After `INSERT_NEW_AT` rounds, dynamic clients are added using **Thompson Sampling** based on trust scores.

3. Model Definitions

- Implements both **SimpleNN** and **SimpleCNN**.
- Chooses CNN for CIFAR-10 (due to image complexity).

4. Training Procedure

- Each round, clients train locally and send updated models.
- Global model is updated by **averaging** local weights.
- Two separate global models:
 - One for **Fed-TS selection**
 - One for **Random selection** (baseline)

5. Drift Detection

- Measures the **weight update distance** (`get_path_drift`) of each client.
- Uses **KMeans clustering** to define a threshold for suspicious updates.
- Updates **Beta distributions** (success/failure counts) to influence future client selection.

6. Visualization

- Plots accuracy over time for both Fed-TS and Random.
- Visualizes **per-client accuracy** trend for deeper insight.

Technologies Used

Component	Tool/Library
Deep Learning	PyTorch (<code>torch</code>)
Dataset	CIFAR-10 / MNIST
Federated Partitioning	Custom + <code>Subset</code>
Visualization	<code>matplotlib</code>
Sampling Algorithm	Thompson Sampling
Clustering	<code>sklearn.cluster.KMeans</code>
Client Behavior Attack	Label-Flipping Strategy
Persistence (Optional)	<code>torch.save()</code>

Key Configuration

Parameter	Value
<code>ROUND_TOTAL</code>	20
<code>CLIENTS_ORIGINAL</code>	5

CLIENTS_DYNAMIC	5
USE_CIFAR	True
ATTACK_LABEL_FLIPPI NG	True
SAVE_MODEL	True (optional)

Why Thompson Sampling?

Problem: Malicious or unproductive clients degrade model performance if selected blindly.

Solution: Use **Thompson Sampling**, a **Bayesian multi-armed bandit** approach that:

- Scores each client based on success/failure (drift)
 - Samples from Beta distribution to **balance exploration vs exploitation**
 - Reduces participation of unreliable clients
-

What Did We Learn?

- **Fed-TS significantly outperforms Random** selection in the presence of attackers.
- Thompson Sampling enables **adaptive trust-based selection**.
- The system is **scalable and dataset-agnostic**, supporting CIFAR-10 and MNIST.
- Visualizations clearly demonstrate the **robustness** of trust-aware client filtering.